

# Leak and Seek

# A Go Runtime Mystery

Adi Abramovitch, Amit Yahav

# Adi Abramovitch



Software Engineer

# Amit Yahav

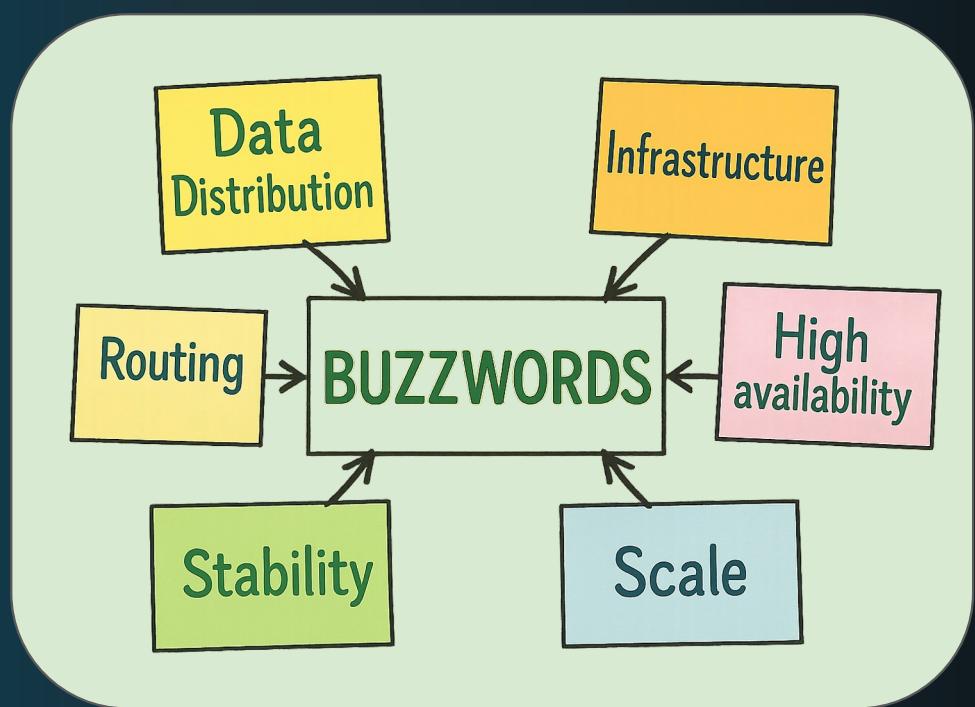


Software Engineer



- Cyber Cyber Cyber
- Simplified Secure Remote Access
- Connecting users to critical infrastructure
- Identity-centric and zero-trust architecture

# The Platform Team



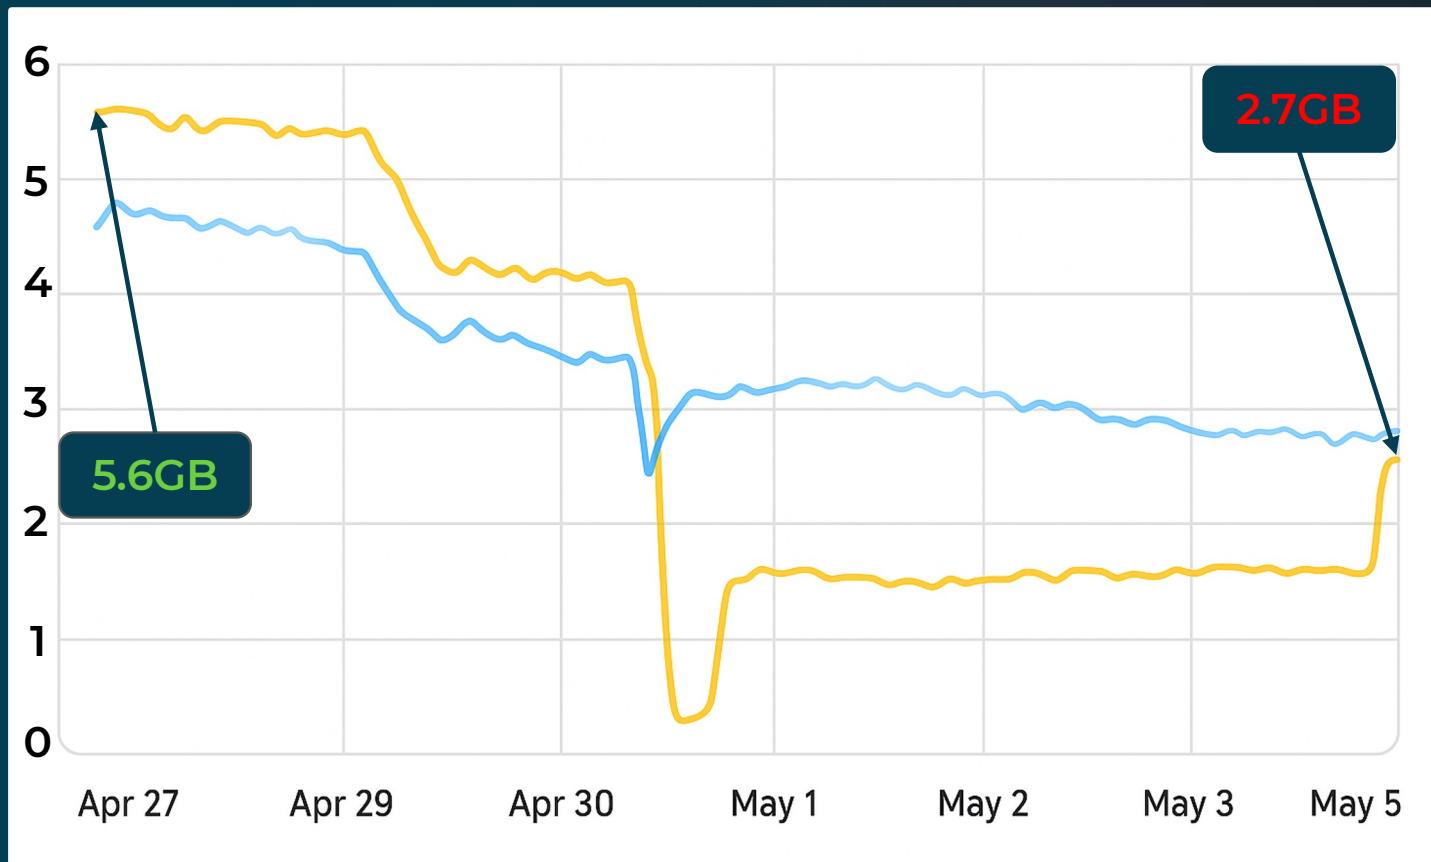
# Agenda

1. The **challenge**
2. Common **debugging** techniques
3. How we identified the **root cause**
4. Our **contribution** to Go and our **open-source** solution

# Chapter One: The Alarm Bells



# Our Free Memory is Decreasing



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# What Is a Memory Leak?



“An unintended loss or escape of something”



but...

Go Is a Garbage-Collected Language



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Potential Cause: A Goroutine Leak

```
○ ○ ○

func main() {
    ch := make(chan byte)

    go func() {
        largeSlice := make([]byte, units.GB)

        // no other routine is pushing data to ch, blocked until the program exists.
        <- ch

        // holds a reference of largeSlice.
        fmt.Printf("slice len: %d", len(largeSlice))
    }()

    // long-running program.
    time.Sleep(60 * time.Minute)
}
```



# How Can Go Help Us Detect Leaks?

# pprof



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# What is PPROF?

A tool to diagnose and visualize performance bottlenecks

Available Profiles:

- CPU
- Block
- Mutex
- Goroutines
- Heap



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Goroutine Profile



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Debug Modes

- **debug=0**

Compressed binary format, used by `go tool pprof` for visualization

- **debug=1**

Human-readable, traces are **grouped** together

- **debug=2**

A full, verbose, human-readable dump of all goroutines



# Goroutine Profile with debug=1

```
○ ○ ○

goroutine profile: total 957
952 @ 0x10413bc58 0x10411f88c 0x104131394c 0x104143f04
    0x104131394b      main.(*EventEmitter).HandleExcessive+0x14b      /Users/user/go/projects/g/
1 @ 0x1040fae5e 0x10413ab04 0x1042ffb74 0x1042ff990 0x1042fcf28 0x104311e14 0x104312728 0x1042db0f7
    0x1042ffb73      runtime/pprof.writeRuntimeProfile+0xb3      /Users/user/go/pkg/mod/golang.org.
    0x1042ff98f      runtime/pprof.writeGoroutine+0x4f      /Users/user/go/pkg/mod/golang.org.
    0x1042fcf27      runtime/pprof.(*Profile).WriteTo+0x147      /Users/user/go/pkg/mod/golang.org.
    0x104311e13      net/http/pprof.handler.ServeHTTP+0x443      /Users/user/go/pkg/mod/golang.org.
    0x104312727      net/http/pprof.Index+0xc7      /Users/user/go/pkg/mod/golang.org.
    0x1042db0f7      net/http.HandlerFunc.ServeHTTP+0x37      /Users/user/go/pkg/mod/golang.org.
    0x1042dcef3      net/http.(*ServeMux).ServeHTTP+0x1b3      /Users/user/go/pkg/mod/golang.org.
    0x1042e3d1b      net/http.serverHandler.ServeHTTP+0xbb      /Users/user/go/pkg/mod/golang.org.
    0x1042d9f7b      net/http.(*conn).serve+0x4fb      /Users/user/go/pkg/mod/golang.org.

1 @ 0x10413bc58 0x1040fe678 0x10413af40 0x104188d18 0x10418982c 0x104213c28 0x10421c8f3 0x10421c8f3 0x1042d503f
    0x10413af3f      internal/poll.runtime_pollWait+0x9f      /Users/user/go/pkg/mod/go:
    0x104188d17      internal/poll.(*pollDesc).wait+0x27      /Users/user/go/pkg/mod/go:
    0x10418982b      internal/poll.(*pollDesc).waitRead+0x1fb      /Users/user/go/pkg/mod/go:
    0x10418981c      internal/poll.(*FD).read+0x1ec      /Users/user/go/pkg/mod/go:
    0x104213c27      net.(*netFD).Read+0x27      /Users/user/go/pkg/mod/go:
    0x10421c8f3      net.(*conn).Read+0x33      /Users/user/go/pkg/mod/go:
    0x1042d503f      net/http.(*connReader).backgroundRead+0x3f      /Users/user/go/pkg/mod/go
```



# Goroutine Profile with debug=1

```
goroutine profile: total 957
952 @ 0x10413bc58 0x10411f88c 0x104131394c 0x104143f04
      0x104131394b          main.(*EventEmitter).HandleExcessive+0x14b

1 @ 0x1040fae5e 0x10413ab04 0x1042ffb74 0x1042ff990 0x1042fcf28 0x10
      0x1042ffb73          runtime/pprof.writeRuntimeProfile+0xb3
      0x1042ff98f          runtime/pprof.writeGoroutine+0x4f
      0x1042fcf27          runtime/pprof.(*Profile).WriteTo+0x147
      0x104311e13          net/http/pprof.handler.ServeHTTP+0x443
```



# Not in our case



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Heap Profile



The Alarm Bells



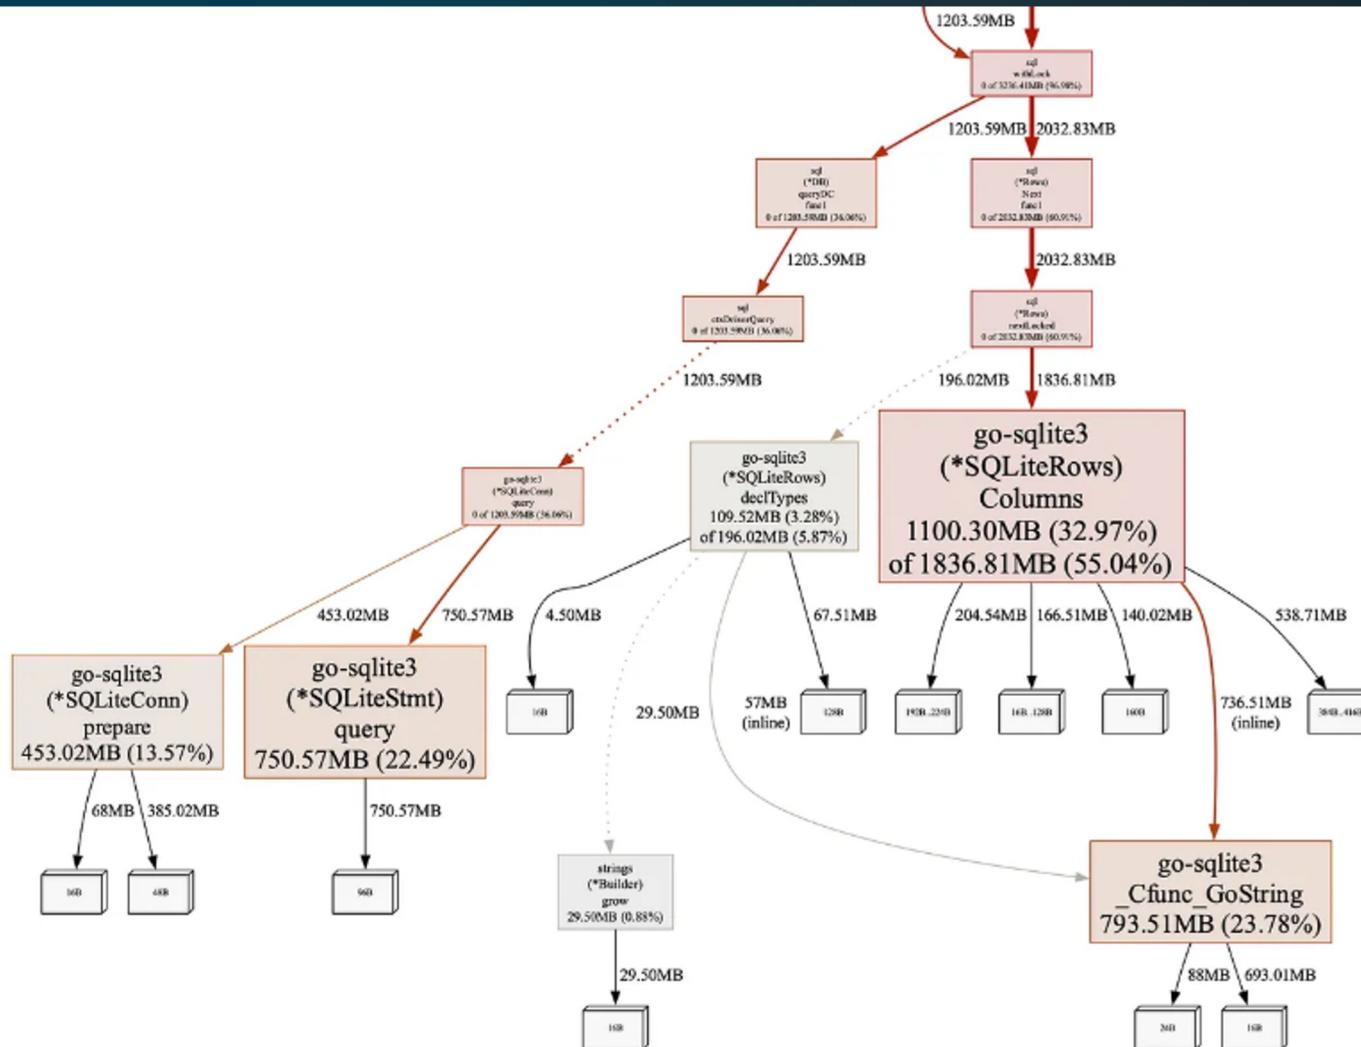
The Three Body Problem



Beyond The Dead-End

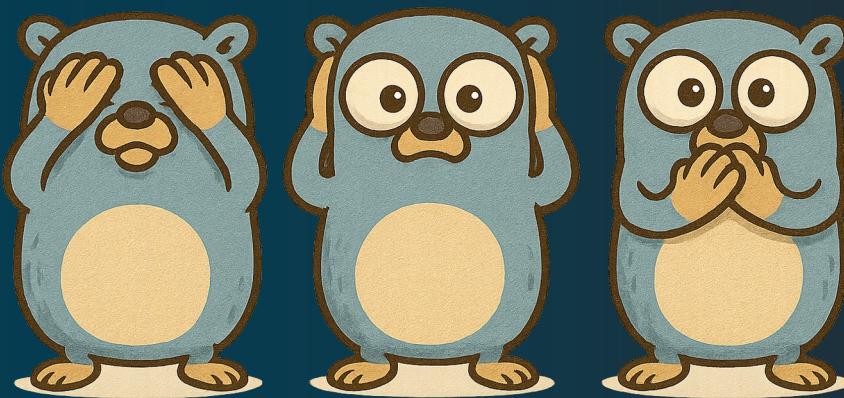


Epilogue



# Chapter Two:

# The Three Body Problem



# SqliteRows

go-sqlite3  
(\*SQLiteRows)  
Columns  
**1100.30MB (32.97%)**



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# SqliteStmt

go-sqlite3  
(\*SQLiteStmt)  
query  
**750.57MB (22.49%)**



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# SqliteConn

go-sqlite3  
(\*SQLiteConn)  
prepare  
**453.2MB (13.57%)**

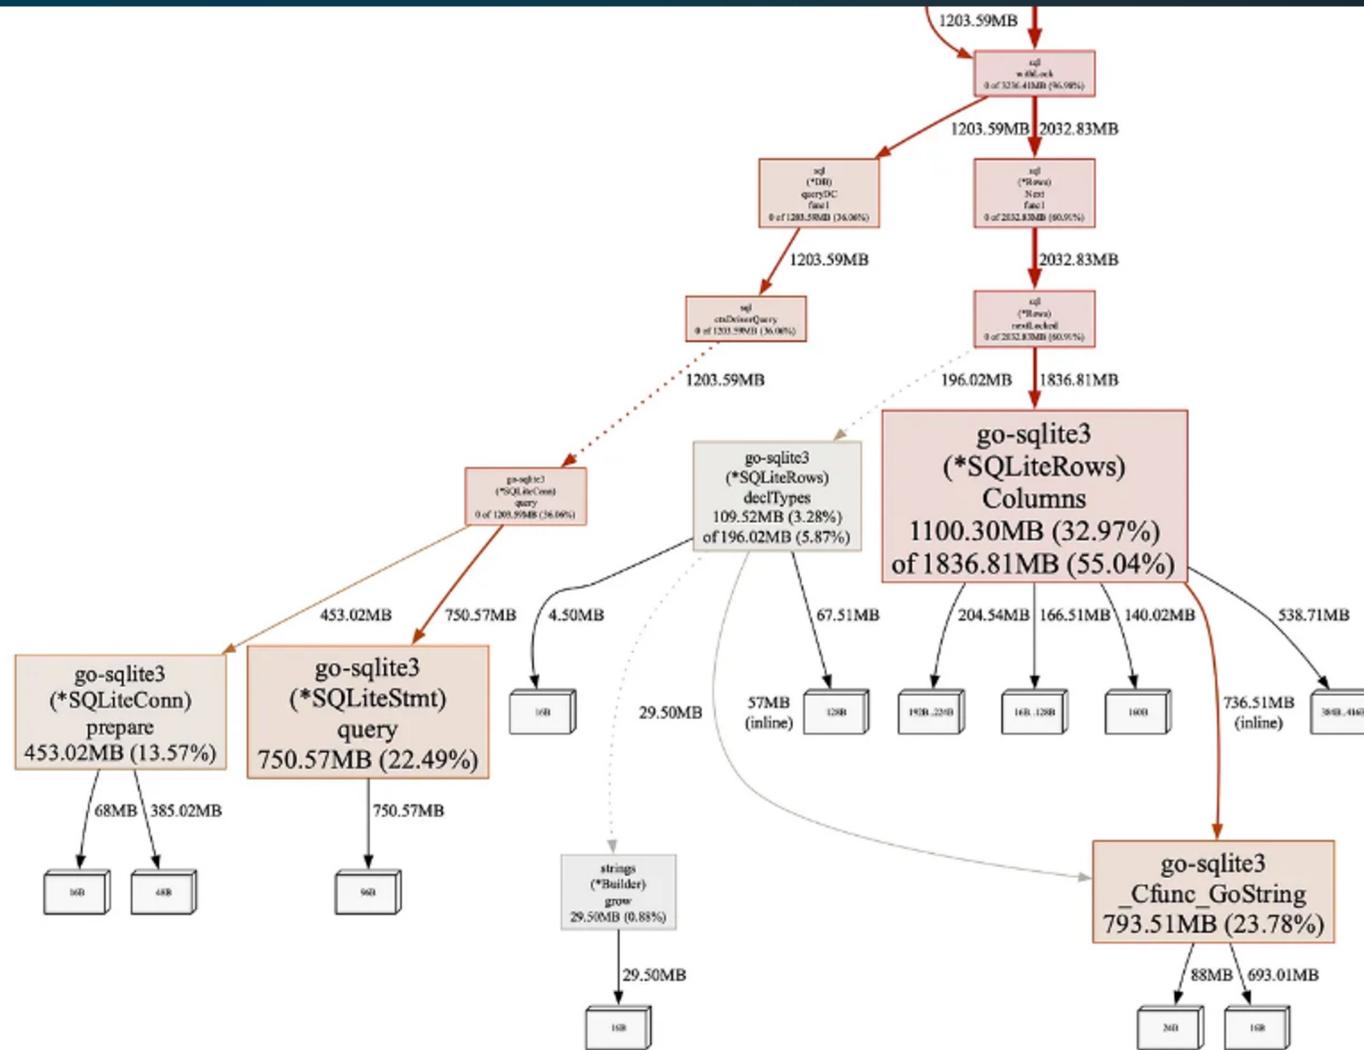


The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue



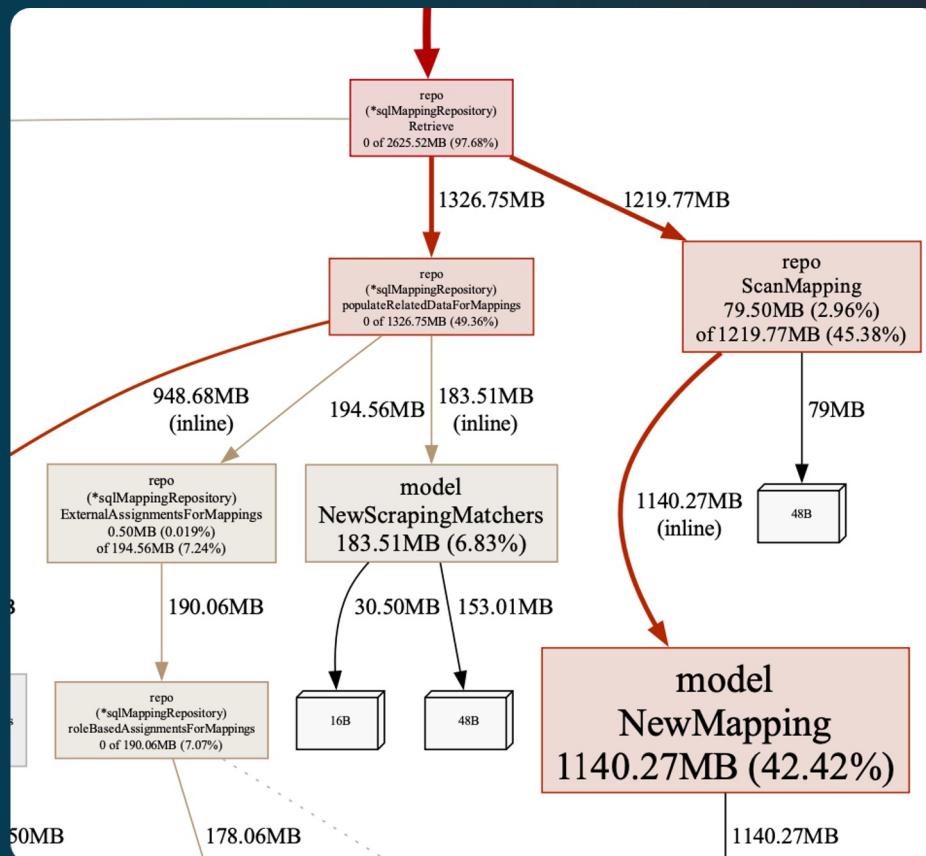
The Alarm Bells

The Three Body Problem

Beyond The Dead-End

Epilogue

# Simulating the Leak





The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Suspect No 1: database/sql

- API for relational databases provided by Go
- Maybe it holds onto the references?
- After thorough inspection: No!



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Suspect No 2: mattn/go-sqlite3 Driver

## Facts:

- SqliteRows → SqliteStmt → SqliteConn
- This library uses CGO
- The version changed
  - Notably: Added **finalizer** to SqliteRows



The Alarm Bells

➤ The Three Body Problem ➤

Beyond The Dead-End ➤

Epilogue

○ ○ ○

```
func (s *SQLiteStmt) query(ctx context.Context, args []driver.NamedValue) (driver.Rows, error)
{
    if err := s.bind(args); err != nil {
        return nil, err
    }

    rows := &SQLiteRows{
        s:          s,
        nc:         int(C.sqlite3_column_count(s.s)),
        cols:       nil,
        decltype:   nil,
        cls:        s.cls,
        closed:     false,
        ctx:        ctx,
    }
    runtime.SetFinalizer(rows, (*SQLiteRows).Close)

    return rows, nil
}
```



# What Is a Finalizer?

A function run by the Go runtime on an object right **before it's garbage collected**

○ ○ ○

```
runtime.SetFinalizer(rows, (*SQLiteRows).Close)
```



# Why Use a Finalizer?

Acts as a **last-resort cleanup** in case the user forgets to call Close() or a similar manual cleanup function



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Finalizers Are Error Prone

Many examples and scenarios are available online

**Sqlite3 got it right!**



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# A Hint in the Docs!



*“A single goroutine runs all finalizers for a program, sequentially. If a finalizer must run for a long time, it should do so by starting a new goroutine.”*

*~runtime/mfinal.go*



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# finq: The Finalizers Queue

```
○ ○ ○
```

```
finq      *finblock // list of finalizers that are to be executed
```

```
○ ○ ○
```

```
func createfing() {
    // start the finalizer goroutine exactly once
    if fingStatus.Load() == fingUninitialized && fingStatus.CompareAndSwap(fingUninitialized, fingCreated) {
        go runfing()
    }
}
```

*~runtime/mfinal.go*



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Maybe runfinq Is Blocked?



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

## POC: A Blocked Finalizer

- Created a finalizer that will block
- **Bingo!**  
All future objects with an associated finalizer will not be released
- **Easy to verify!**  
The blocked finalizer routine can be seen on the goroutine profile



# Goroutine Profile with debug=1

○ ○ ○

```
goroutine profile: total 6

1 @ 0x104561cb8 0x104541e74 0x104541e55 0x1045630d8 0x10457a670 0x104712324 0x10471231d 0x1047122c5 0x104509a84
0x104569a74
#  0x1045630d7 internal/sync.runtime_SemacquireMutex+0x27  /Users/projects/go/pkg/mod/golang.org/toolchain@v0.0.1-
go1.24.1.darwin-arm64/src/runtime/sema.go:95
#  0x10457a66f internal/sync.(*Mutex).lockSlow+0x16f      /Users/projects/go/pkg/mod/golang.org/toolchain@v0.0.1-
go1.24.1.darwin-arm64/src/internal/sync/mutex.go:149
#  0x104712323 internal/sync.(*Mutex).Lock+0xb3        /Users/projects/go/pkg/mod/golang.org/toolchain@v0.0.1-
go1.24.1.darwin-arm64/src/internal/sync/mutex.go:70
#  0x10471231c sync.(*Mutex).Lock+0xac          /Users/projects/go/pkg/mod/golang.org/toolchain@v0.0.1-
go1.24.1.darwin-arm64/src/sync/mutex.go:46
#  0x1047122c4 main.(*Leaky).Close+0x54        /Users/projects/Library/Application
Support/JetBrains/GoLand2025.1/scratches/scratch_76.go:20
#  0x104509a83 runtime.runfinq+0x3d3       /Users/projects/go/pkg/mod/golang.org/toolchain@v0.0.1-
go1.24.1.darwin-arm64/src/runtime/mfinal.go:275
```





# Eureka!



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# No Trace of the finq Routine...



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Conclusion

The blocked finalizer is not the leak's root cause



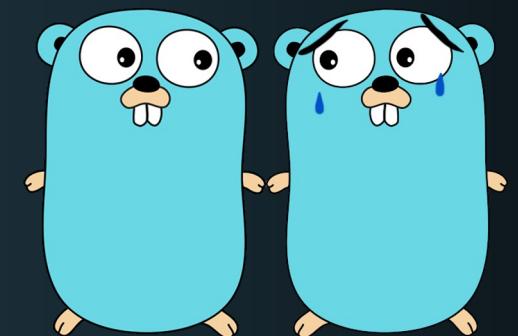
The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Dead End?



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Chapter Three:

# Beyond the Dead-End



Weapons Free: Debug Mode

# Goref

*cloudwego/goref*



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Weapons Free: Debug Mode

- Build our app in debug mode to utilize the symbol table
- Released a new version to the customer
- Obtained core dump
- Generated goref profile using ***grf core exec\_file core\_file***
- Observed the output using ***go tool pprof -http=:5079 ./grf.out***



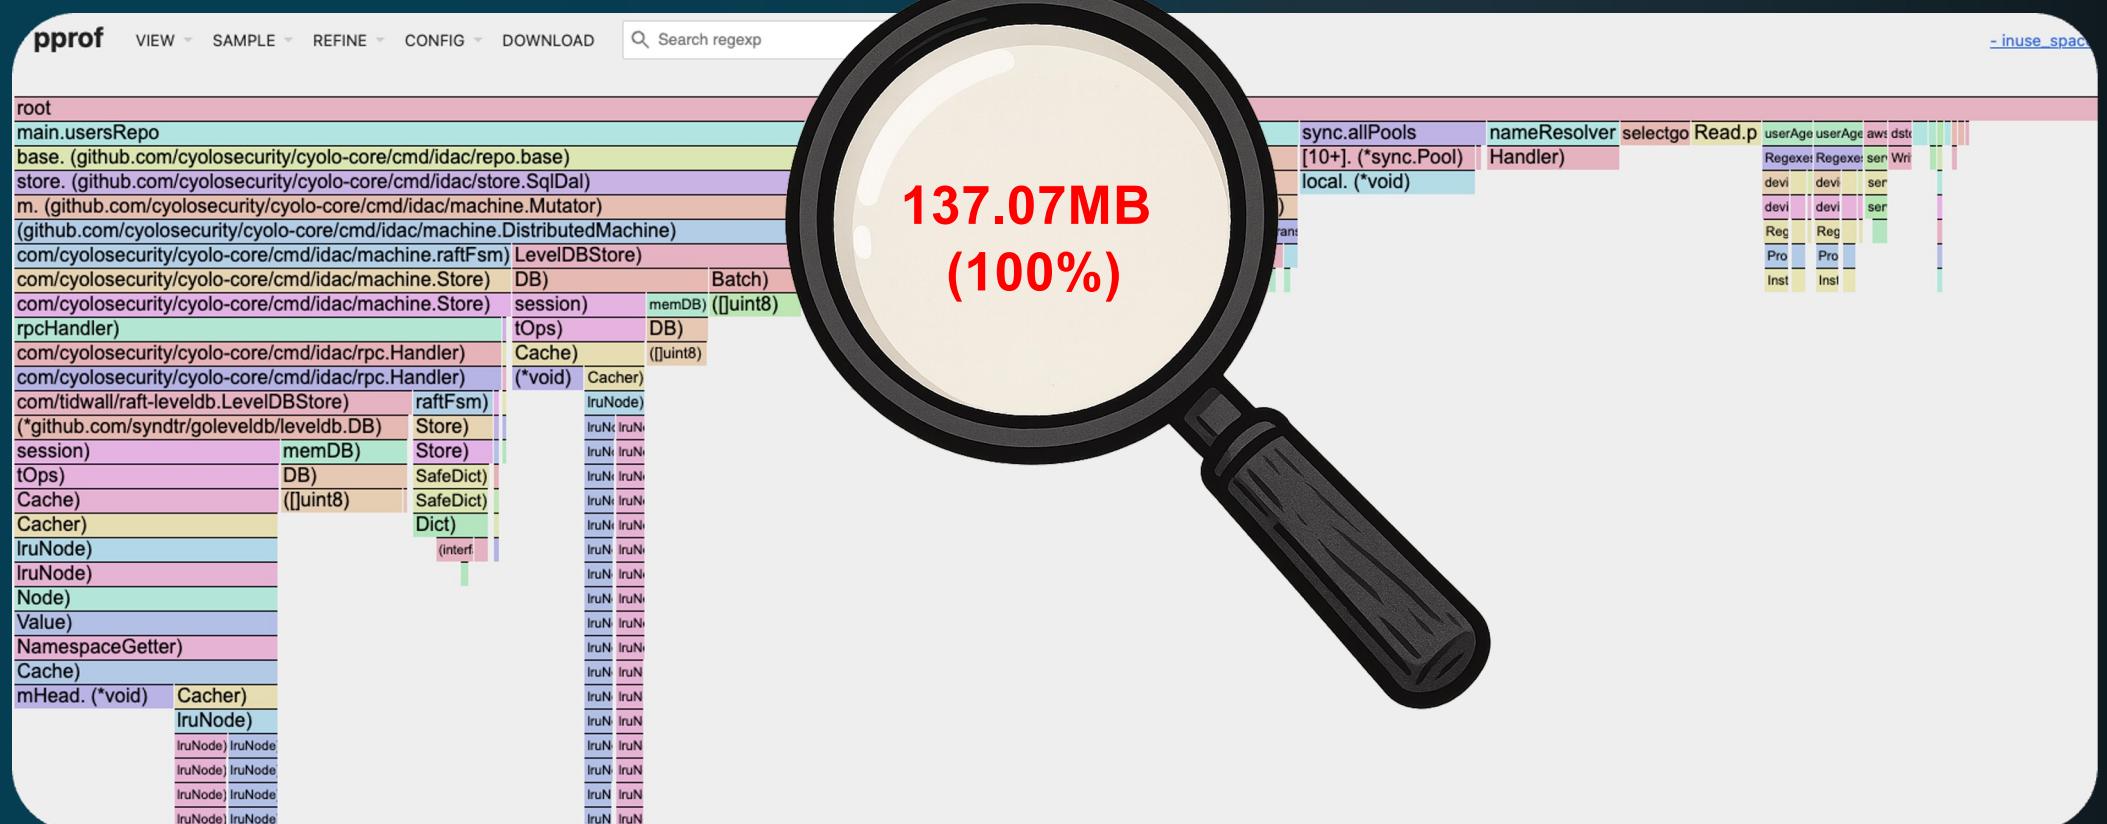
The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

*go tool pprof -http=:5079 ./grf.out*



The Alarm Bells

➤ The Three Body Problem ➤

Beyond The Dead-End

➤ Epilogue

# What Do We Know So Far

- We identified the **flows** that lead to the leaking allocation
- We identified an additional **finalizer** in sqlite3
- We recreated a leak using a **blocked** finalizer
- We retrieved a goroutine profile from the **customer**
- **No trace** for runfinq in the goroutine profile



# Conclusion, yet again...

The blocked finalizer is not the leak's root cause



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Doubting Ourselves

- A Go runtime bug?
- Goroutine profile is not accurate?
- A blocked finalizer?



The Alarm Bells

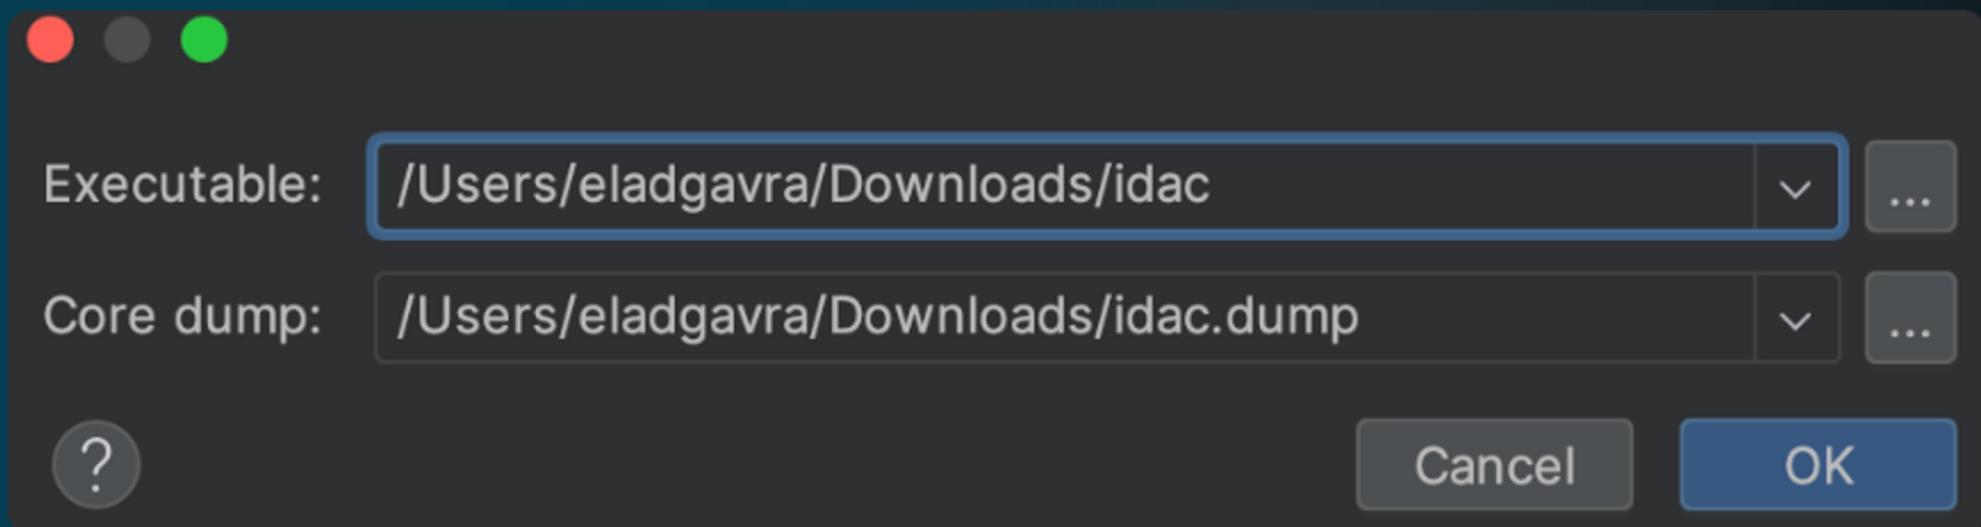
➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Analyzing the Core Dump

- GoLand > Run > Open Core Dump



# Our Next Objectives

- Trace of runfinq's routine
- What is it currently doing?
- Is it blocked?



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# The Infamous: Goroutine 34

The screenshot shows a debugger interface with two main panes. The left pane displays a stack trace for Goroutine 34, which is currently blocked in the `runtime.gopark` function. Several lines in the stack trace are highlighted with red boxes: `runtime.gopark`, `runtime.selectgo`, `smb2.(*File).close`, `runtime.call32`, `runtime.runfinq`, and `runtime.goexit`. The right pane shows the value of the `gp.waitsince` variable, which is a struct containing fields like `result`, `cas0`, `order0`, `nsends`, `block`, `cas1`, `order1`, `gp`, `lockorder`, `t0`, and `norder`.

```
gp.waitsince
  result = {int64} 2719385194558727
  cas0 = {*runtime.scase | 0xc00c352cf0}
  order0 = {uint16 | 0xc00c352ca0} 1
  nsends = {int} 0
  block = {bool} true
  cas1 = {[65536]runtime.scase | 0xc00c352cf0} len:65536
  order1 = {[131072]uint16 | 0xc00c352ca0} len:131072
  gp = {*runtime.g | 0xc00019a000}
  lockorder = {[]uint16} len:2, cap:0
    t0 = {int64} 0
    norder = {int} 2
```



# The Infamous: Goroutine 34

**gp.waitsince** - the nanosecond timestamp of a goroutine's entry into a waiting state (parked).

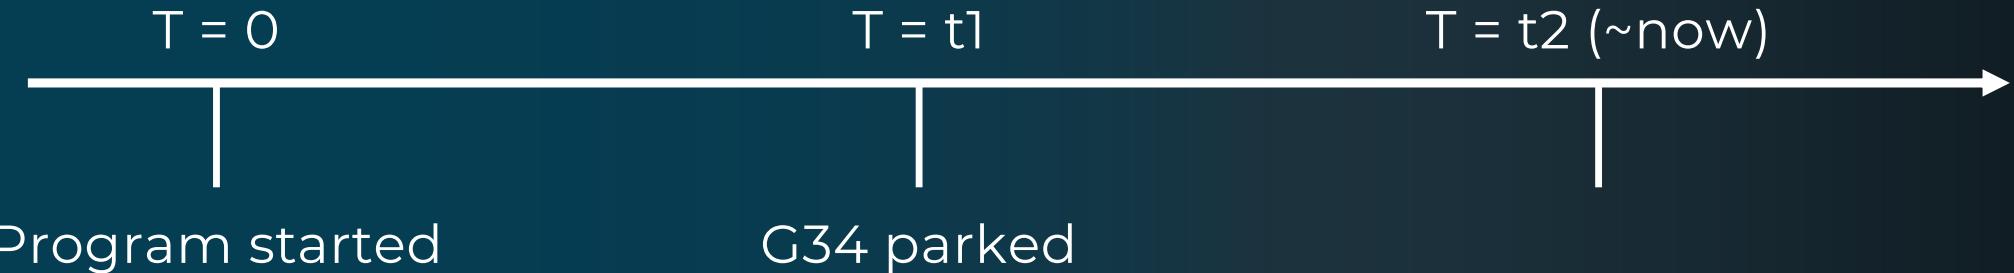
The screenshot shows the GDB interface with the 'Threads & Variables' tab selected. On the left, the stack trace for Goroutine 34 is visible, showing calls from runtime.gopark down to runtime.createfinq. On the right, the variable 'gp.waitsince' is highlighted with a red box. Its value is listed as 'result = {int64} 2719385194558727'. Below it, other variables are listed: cas0, order0, nsends, block, cas1, order1, gp, lockorder, t0, and norder.

```
gp.waitsince
  result = {int64} 2719385194558727
  cas0 = {*runtime.scase | 0xc00c352cf0}
  order0 = {uint16} 1
  nsends = {int} 0
  block = {bool} true
  cas1 = {*[65536]runtime.scase | 0xc00c352cf0} len:65536
  order1 = {*[131072]uint16 | 0xc00c352ca0} len:131072
  gp = {*runtime.g | 0xc00019a000}
  lockorder = {[]uint16} len:2, cap:0
    t0 = {int64} 0
    norder = {int} 2
```



# How long has G34 been parked?

$$t_2 - t_1 = ?$$



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# How long has G34 been parked?

$$t_2 - t_1 = ?$$



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# How long has G34 been parked?

$$t_2 - t_1 = ?$$



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# How long has G34 been parked?

$$t_2 - t_1 = 197 \text{ minutes}$$



The Alarm Bells

➤ The Three Body Problem

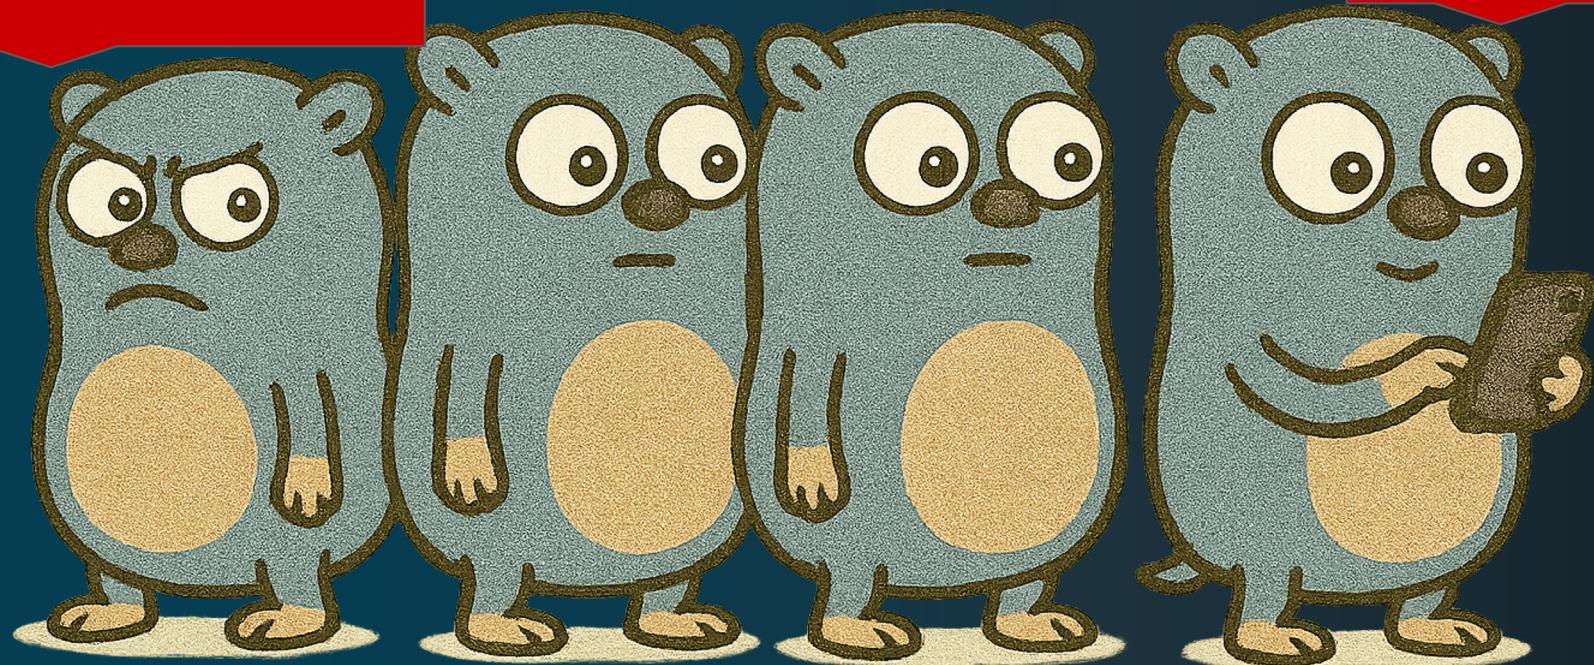
➤ Beyond The Dead-End

➤ Epilogue

# Finalizers Queue

(\*SqliteRows).Close

smb2.(\*File).close



The Alarm Bells

The Three Body Problem

Beyond The Dead-End

Epilogue

# hirochachacha/go-smb2

```
○ ○ ○

func (fs *Share) newFile(r CreateResponseDecoder, name string) *File {
    fd := r.FileId().Decode()

    fileStat := &FileStat{
        CreationTime: time.Unix(0, r.CreationTime().Nanoseconds()),
        LastAccessTime: time.Unix(0, r.LastAccessTime().Nanoseconds()),
        LastWriteTime: time.Unix(0, r.LastWriteTime().Nanoseconds()),
        ChangeTime: time.Unix(0, r.ChangeTime().Nanoseconds()),
        EndOfFile: r.EndOfFile(),
        AllocationSize: r.AllocationSize(),
        FileAttributes: r.FileAttributes(),
        FileName: base(name),
    }

    f := &File{fs: fs, fd: fd, name: name, fileStat: fileStat}
    runtime.SetFinalizer(f, (*File).close)
    return f
}
```



# Issues Identified

- `(*File).close` did not run in a separate goroutine
- **197 minutes** for an I/O operation is not reasonable
- SMB finalizer is blocked



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Why SMB Finalizer is Blocked?

Race condition in **(\*File).close** method



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# The Race

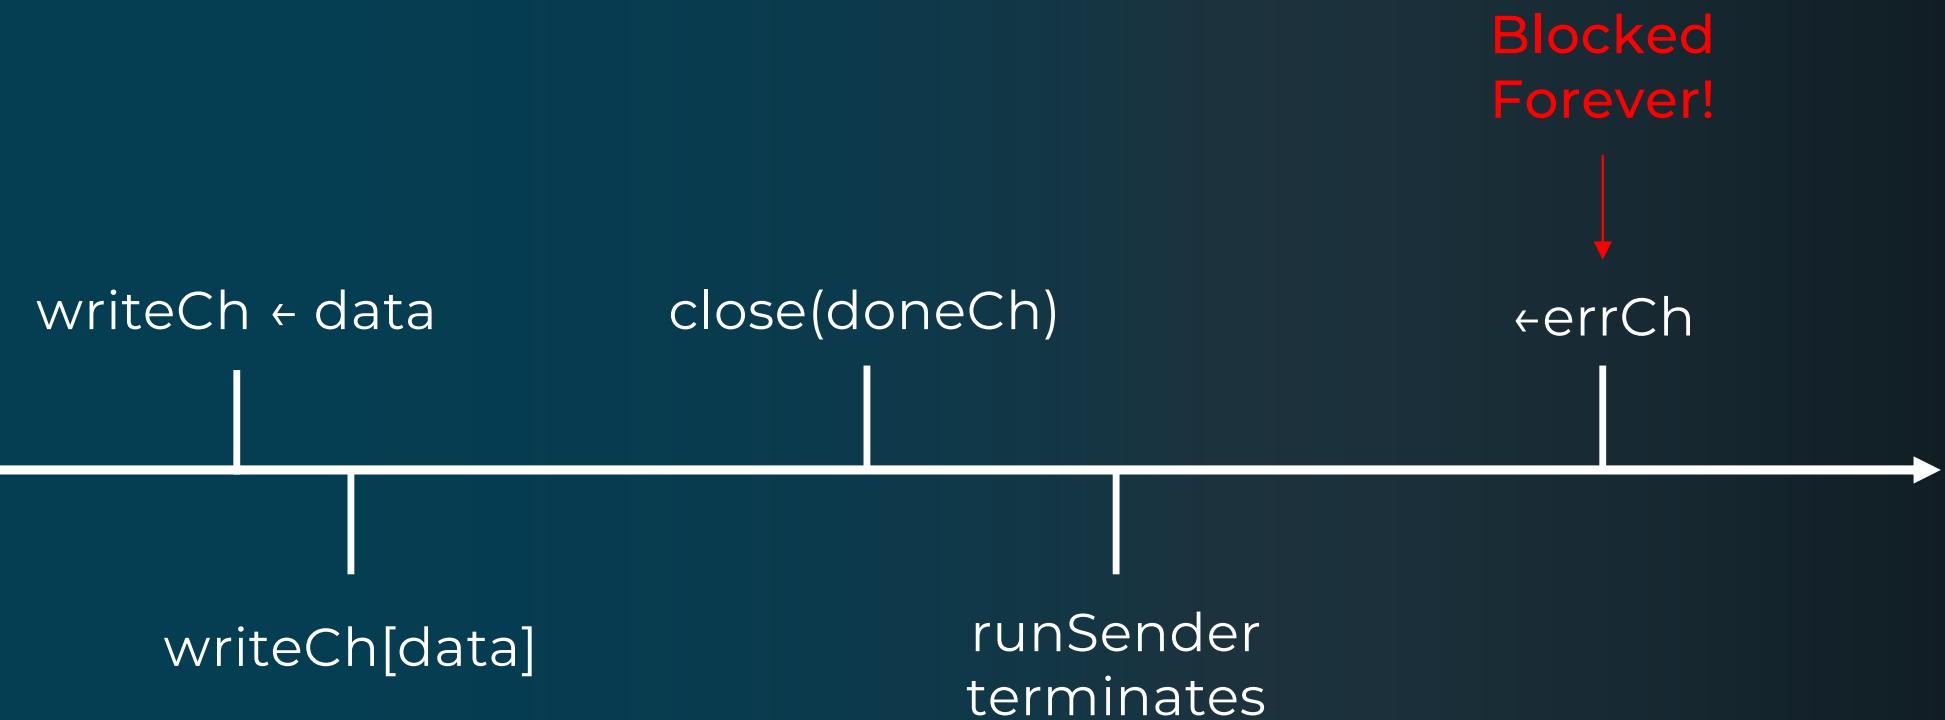
Finalizer: (\*File).close

```
○ ○ ○  
...  
select {  
    case conn.writeCh <- rr.pkt:  
        select {  
            case err = <-conn.errCh:  
                ...  
            case <-ctx.Done():  
                return  
        }  
    case <-ctx.Done():  
        ...  
}
```

```
○ ○ ○  
func (conn *conn) runSender() {  
    for {  
        select {  
            case <-conn.doneCh:  
                return  
            case pkt := <-conn.writeCh:  
                _, err := conn.t.Write(pkt)  
                conn.errCh <- err  
        }  
    }  
}
```



# The Race



# The Race

Finalizer: (\*File).close

```
○ ○ ○  
...  
select {  
    case conn.writeCh <- rr.pkt:  
        select {  
            case err = <-conn.errCh:  
                ...  
            case <-ctx.Done():  
                return  
        }  
    case <-ctx.Done():  
        ...  
}
```

```
○ ○ ○  
func (conn *conn) runSender() {  
    for {  
        select {  
            case <-conn.doneCh:  
                return  
            case pkt := <-conn.writeCh:  
                _, err := conn.t.Write(pkt)  
                conn.errCh <- err  
        }  
    }  
}
```



# The Fix: Bypassing The Race

○ ○ ○

```
ctx, cancel := context.WithCancel(parentCtx)
defer cancel()
...
```



# The Race

Finalizer: (\*File).close

```
○ ○ ○  
...  
select {  
    case conn.writeCh <- rr.pkt:  
        select {  
            case err = <-conn.errCh:  
                ...  
            case <-ctx.Done():  
                return  
        }  
    case <-ctx.Done():  
        ...  
}
```

```
○ ○ ○  
func (conn *conn) runSender() {  
    for {  
        select {  
            case <-conn.doneCh:  
                return  
            case pkt := <-conn.writeCh:  
                _, err := conn.t.Write(pkt)  
                conn.errCh <- err  
        }  
    }  
}
```



# Solved!



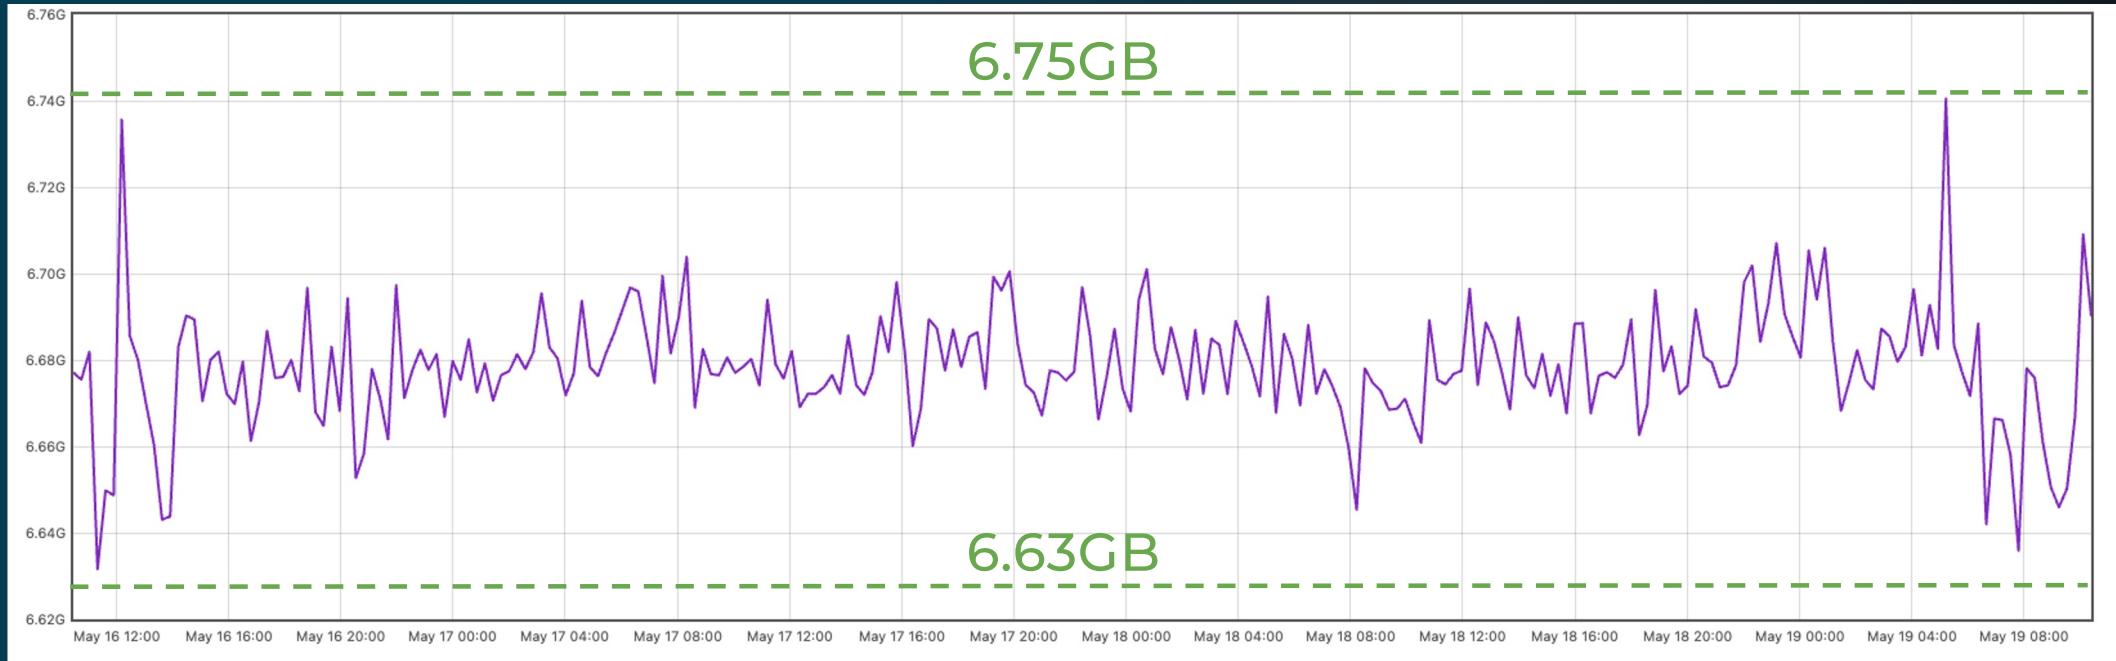
The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Free Memory



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# The Implication

Package A

Blocking Finalizer

Package B

Memory Leak



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue



Epilogue:  
**From Bug to Better Go**

# Open Questions

- Why Goref did not capture our leak?
- Why was runfinq missing from the goroutine profile?



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Why Goref Did not Capture Our Leak?

○ ○ ○

```
// finblock is allocated from non-GC'd memory, so any heap pointers
// must be specially handled. GC currently assumes that the finalizer
// queue does not grow during marking (but it can shrink).
type finblock struct {
    _           sys.NotInHeap
    alllink    *finblock
    next       *finblock
    cnt        uint32
    _           int32
    fin        [(_FinBlockSize - 2*goarch.PtrSize - 2*4) / unsafe.Sizeof(finalizer{})]finalizer
}
```



# Why Was runfinq Missing From The Goroutine Profile?

- We recreated the blocked routine scenario
- ‘runfinq’ is visible with **debug=1**
- ‘runfinq’ is **not** visible with **debug=2**

A bug!



## Follow-up Go Issues

- **[#73011]** Goroutine profile at debug=2 hides the start of the finalizer goroutine. Fixed in Go 1.25!
- **[#72950]** Explicit documentation for blocked finalizers and their consequences
- **[#72948]** A metric for detecting slow or stalled finalizer queues
- **[#72949]** A dedicated debug mode for finalizers



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

Are we done? Not yet!

We need to detect the next blocking finalizer,  
**effortlessly**, starting **yesterday**.



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Implementing a Proactive Monitoring Solution

1. Generate a **disposable** object with an associated finalizer
2. Measure the **duration** from enqueueing until execution
3. Export this data as a **metric**
4. Upon blockage: **notify** with a trace of the blocking finalizer



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Metric: Finalizer Duration



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Introducing: finq

A tool for monitoring Go's finalizer routine responsiveness.

```
○ ○ ○

go finq.Monitor(ctx, &MonitorOpts{
    StallingInterval: time.Minute,
    OnStalling: func(d time.Duration) {
        log.Println("waiting for finalizer to execute for:", d)

        // you can use this functionality to retrieve the stacktrace of the blocking routine
        trace := finq.GetFinalizerStackTrace()
        if trace != "" {
            log.Println("blocking routine trace:", trace)
        }
    },
    OnComplete: func(d time.Duration) {
        log.Println("finalizer executed after:", d)
    },
    ImmediatelyTriggerGC: false,
})
```

[github.com/cyolosecurity/finq](https://github.com/cyolosecurity/finq)



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# AddCleanup in Go 1.24

○ ○ ○

```
cleanupFunc := func(fileName string) {
    if err := os.Remove(fileName); err == nil {
        fmt.Println("temp file has been removed")
    }
}

runtime.AddCleanup(&tempFile, cleanupFunc, tempFile.Name())
```



## AddCleanup in Go 1.24

AddCleanup is much less error-prone than SetFinalizer

Would it have solved our problem?

No



The Alarm Bells



The Three Body Problem



Beyond The Dead-End



Epilogue

# Summary

Unusual Leak

Debugging  
Arsenal

Contribution



The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End



➤ Epilogue

# Our Key Takeaways

Persistency

Collaboration

Solidarity



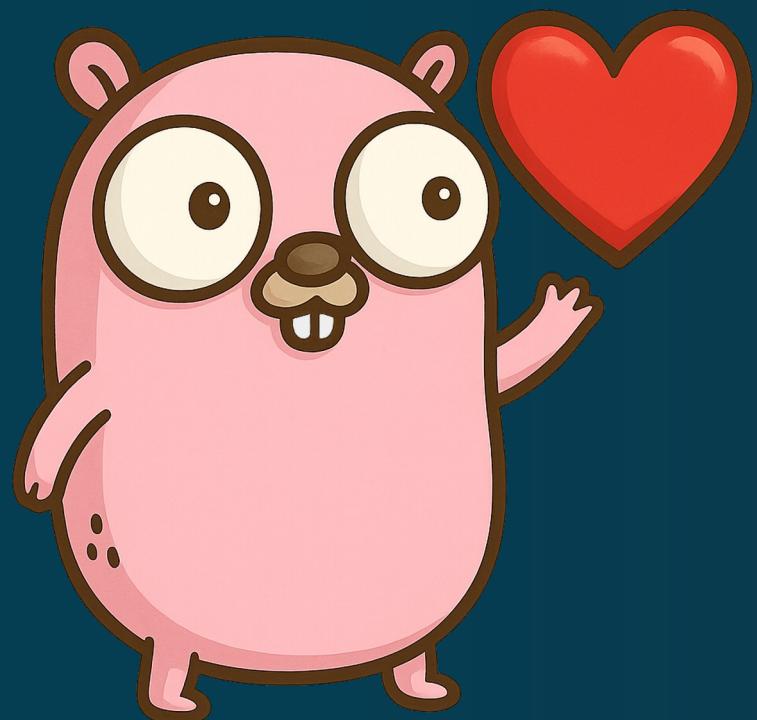
The Alarm Bells

➤ The Three Body Problem

➤ Beyond The Dead-End

➤ Epilogue

# Thank You!



## Questions?

